

Final Year Project Report

Final Unit – Full Project

**AI-Agent for Monitoring and Improving Student Academic
Performance**

Salim Mohamed Hussein

Hanad Hayir Mohamud

A report Submitted in partial fulfilment of the degree of

BSc in Information Technology

Supervisor: Mohamed Abdullahi Ali



Department of **Information Technology**

Faculty of computing, SIMAD university

April 2026

DECLARATION. A

“We declare that the following is our own work and does not contain any unacknowledged work from any other sources. This project was undertaken to fulfil the requirements of the bachelor’s degree program in Information Technology at SIMAD University”.

Name:

Student Signature:

Name:

Student Signature:

DATE:

DECLARATION. B

I hereby declare that I have reviewed this project titled AI-Agent for Monitoring and Improving Student Academic Performance and, in my opinion, it is adequate in terms of scope and quality for consideration as a project.

Name of Supervisor: Mohamed Abdullahi Ali

Signature:

Date:

ACKNOWLEDGEMENTS

Student One:

This project report could not have been prepared without the help and encouragement of various people. Hence, I would like to extend my gratitude to my guide, mentor, and of course, my supervisor, Mohamed Abdullahi Ali. His support provided me with the proper guidelines needed to successfully complete this project. I would also like to thank everyone who contributed to this work, as without their assistance, this project would have been impossible to accomplish. The information gathered from websites, books, journals, and other sources has been instrumental in shaping this project, and I deeply appreciate all those who contributed to its success

Student Two:

I extend heartfelt gratitude to my supervisor Mohamed Abdullahi Ali, a candidate at SIMAD University, for his invaluable guidance and precise feedback. The Faculty of Computing at SIMAD University provided essential resources and support, which significantly enhanced the quality of this project. Their commitment to academic excellence through workshops has been instrumental in refining my skills. In closing, I sincerely appreciate all those who have contributed to making this project, " AI-Agent for Monitoring and Improving Student Academic Performance" a reality. Their collective efforts have not only shaped the success of this project but have also enriched the academic discourse to which it belongs.

ABSTRACT

Student dropout and academic failure remain persistent challenges in higher education, yet most institutions still rely on reactive, end-of-semester reporting that identifies struggling students only after intervention is no longer effective. This project develops and evaluates an intelligent, event-driven agent that continuously monitors student academic performance and produces personalized, early intervention messages while a semester is still in progress.

The system integrates three core components. A Long Short-Term Memory (LSTM) neural network, built with TensorFlow and Keras, processes longitudinal behavioral sequences including engagement, assessment scores, and submission timing to output a per-subject risk score between 0.0 and 1.0. A reasoning engine powered by a Large Language Model then transforms these risk scores and behavioral signals into empathetic, subject-specific intervention messages. A Python-based Agent Controller orchestrates the pipeline, while a MySQL database stores performance sequences and a React dashboard allows advising staff to review and dispatch interventions, with a dedicated student portal for receiving them.

The model was trained and validated on the Open University Learning Analytics Dataset (OULAD), comprising 28,297 student-subject enrollments after preprocessing. Evaluated on a held-out test set ($N = 4,238$) using a student-level split to prevent data leakage, the LSTM achieved a Test AUC of 0.9839 and an overall classification accuracy of 94% when given a student’s full behavioural sequence. At the week-13 early-warning checkpoint, where only the first part of the semester is observed, performance is necessarily lower but still informative (AUC 0.84), and a milestone analysis shows discrimination rising steadily from an AUC of 0.70 at day 30 to 0.98 over the full term. A benchmark against CNN, Random Forest, and Logistic Regression baselines on the same split confirmed that the LSTM is the strongest model. These results confirm that the combination of sequential deep learning and language-model reasoning provides a scalable foundation for proactive academic support, with a clear adaptation path toward in-person universities through attendance and learning-management data.

Keywords: Long Short-Term Memory (LSTM), Student Performance Prediction, Early Warning System, Large Language Models, Learning Analytics, Agentic AI.

TABLE OF CONTENTS

DECLARATION. A.....	ii
DECLARATION. B.....	iii
ACKNOWLEDGEMENTS	iv
ABSTRACT	v
TABLE OF CONTENTS	vii
LIST OF TABLES	ix
LIST OF FIGURES	x
CHAPTER ONE INTRODUCTION	11
1. Introduction.....	11
1.1. Background of the study.....	12
1.2. Problem Statement	13
1.3. Research Objectives.....	15
1.4. Scope of the study	15
1.5. Significance of the Study.....	16
1.6. Definition of terms	17
CHAPTER TWO: LITERATURE REVIEW	18
2. Introduction to the Chapter	18
2.1. Educational Data Mining (EDM) and Predictive Student Analytics	19
2.2. Deep Learning and Longitudinal Time-Series Forecasting.....	20
2.3. Generative AI and Automated Pedagogical Mentorship	21
2.4. Ethical Considerations in AI-Driven Learning Analytics.....	21
2.5. Identification of the Research Gap.....	22
CHAPTER 3: RESEARCH METHODOLOGY	23
3. Research Design	23

3.1. System Development Methodology (CRISP-DM)	23
3.2. Data Collection and Dataset Analysis.....	24
3.3. System Requirements Analysis	25
3.4. System Design and Architecture	26
3.5. Development Tools and Technologies	27
3.6. Testing and Evaluation.....	27
3.7. Ethical Considerations	27
3.8. Schedule Feasibility	28
CHAPTER FOUR: RESULTS AND DISCUSSIONS	29
4. Introduction.....	29
4.1. Architectural Convergence Analysis	29
4.2. Model Performance and Validation on the Test Set	29
4.3. Discussion of the Precision-Recall Asymmetry	31
4.4. Early-Warning Milestone Analysis	31
4.5. System Verification and Dashboard Interface Analysis.....	32
4.6. Comparison with Existing Systems	35
4.7. Comparative Model Evaluation	36
CHAPTER FIVE: CONCLUSIONS AND RECOMMENDATIONS	38
5. Introduction.....	38
5.1. Summary of Empirical Findings	38
5.2. Conclusions	39
5.3. Recommendations	39
5.4. Limitations and Future Work.....	41
Reference.....	42

LIST OF TABLES

Table 1.1 Operational Definition of Terms	17
Table 2.1 Critical Comparison of Predictive Algorithms in Learning Analytics	20
Table 3.1 Project Implementation Timeline	28
Table 4.1 Test Set Confusion Matrix Boundaries ($N = 4,238$)	30
Table 4.2 Comprehensive Test Set Classification Report	30
Table 4.3 Longitudinal Milestone Analysis Performance Over Time	31
Table 4.4 Comparison of the Proposed System with Existing Early-Warning Systems	36
Table 4.5 Comparative Performance of the LSTM, CNN, and Classical Baselines	37

LIST OF FIGURES

Figure 3.1 system design and architecture	26
Figure 4.1 Institutional Advisor Analytics Console and Cohort Risk Manifest]	33
Figure 4.2 Student #6516 Single Profile Advisor Analysis Screen	34
Figure 4.3 Student Dashboard.....	35
Figure 4.4 Model comparison across evaluation metrics on the held-out test set	37

CHAPTER ONE INTRODUCTION

1. Introduction

Education is widely recognized as one of the most powerful drivers of individual opportunity and national development. As universities expand access and enrolment numbers grow, the ability of institutions to support each student effectively becomes increasingly difficult. A significant proportion of students who enrol in higher education either withdraw before completion or fail to meet the academic standards required to progress, and student dropout remains one of the most complex and costly challenges facing education systems worldwide (Niang et al., 2026). Identifying these students early, while corrective support can still make a difference, has therefore become a central concern in modern educational management.

Traditionally, academic monitoring has depended on periodic examinations and end-of-semester result reviews. While these methods provide a formal record of performance, they are inherently retrospective: they describe what has already happened rather than anticipating what is likely to happen next. By the time a failing grade is recorded, the opportunity to redirect a student through tutoring, counselling, or workload adjustment has frequently passed. This reactive posture is a structural limitation of conventional monitoring, and it is precisely the gap that data-driven, predictive systems are designed to address.

Advances in machine learning, and particularly in deep learning architectures designed for sequential data, now make it possible to model a student's academic trajectory as it unfolds. Recurrent neural networks such as the Long Short-Term Memory (LSTM) network, first introduced to overcome the limitations of earlier recurrent models in learning long-term dependencies (Hochreiter & Schmidhuber, 1997), can learn from the temporal patterns hidden in a student's week-by-week behaviour, including their engagement with learning materials, the timing of their submissions, and the trend of their assessment scores. When combined with a reasoning layer capable of generating natural, human-readable guidance, these models can move beyond simple prediction to deliver actionable, personalized support. This project applies these technologies to design an intelligent agent that monitors student academic performance continuously and enhances it through timely, empathetic intervention.

1.1. Background of the study

The challenge of student retention and academic success is a global concern that affects institutions across both developed and developing contexts. Universities invest substantial resources in admitting and teaching students, yet a meaningful share of those students do not graduate on schedule, and some do not graduate at all. The cost of this attrition is borne not only by institutions, which lose tuition revenue and see their performance indicators decline, but also by the students themselves, who may carry financial and personal burdens without obtaining the qualification they pursued.

Over the past two decades, the field of learning analytics has emerged to address this problem by applying statistical and computational techniques to the large volumes of data that educational systems generate. Learning management systems, student information systems, and assessment platforms now capture detailed records of how students interact with their courses. These records, when analysed effectively, contain early signals of disengagement and difficulty long before they manifest as a failing grade. The central premise of learning analytics is that these signals can be detected, modelled, and acted upon in time to change outcomes (Romero & Ventura, 2020).

Early warning systems built on this premise have been deployed at several large universities, most notably through the Course Signals initiative at Purdue University, which used predictive modelling to flag at-risk students with a traffic-light indicator and prompt instructor outreach (Arnold & Pistilli, 2012). These systems demonstrated that predictive analytics can improve retention, but they also revealed important limitations. Many relied on relatively simple statistical models, such as logistic regression, that captured correlations rather than the temporal dynamics of student behaviour, and most stopped at the point of generating a risk flag, leaving the design of the actual intervention entirely to human staff. As student populations grow, the manual effort required to craft individualized guidance for every flagged student becomes a serious bottleneck.

More recent research has shown that deep learning models, and LSTM networks in particular, can improve on these earlier approaches by learning directly from the temporal sequence of student interactions. Studies using clickstream and learning-management-system logs have demonstrated that LSTM-based models consistently outperform conventional baselines such as logistic regression, random forest, and gradient boosting when predicting dropout and course outcomes from early behavioural data (Santos & Henriques, 2024; Ghazvini et al., 2024). This

body of work establishes the sequential modelling of student behaviour as a credible foundation for early warning.

The recent emergence of large language models has opened a new possibility for closing the gap between detection and response. These models can generate fluent, context-sensitive text, and a growing body of research indicates that they can provide detailed, personalized feedback to students that is often comparable in quality to human-authored feedback (Dai et al., 2023). This means that the step of translating a numerical risk score into a personalized, supportive message can itself be automated. By pairing a sequential prediction model that identifies who is at risk with a language model that explains what each student might do about it, an institution can deliver individualized academic support at a scale that would be impossible through manual effort alone. This project is situated at the intersection of these two developments, combining sequential deep learning with language-model reasoning inside a single autonomous agent.

It is also important to acknowledge the context in which such a system would be deployed. Much of the publicly available data for academic prediction originates from fully online or distance-learning institutions, where every interaction is digitally logged. The Open University Learning Analytics Dataset, which underpins this study, is one such resource, providing demographic data together with aggregated records of student interactions in a virtual learning environment (Kuzilek et al., 2017). In-person universities, including many across the region, capture engagement differently, often through attendance and a more limited set of digital traces. A practical system must therefore be designed so that its core predictive engine remains valid even as the specific source of engagement data changes, an adaptability that is treated as an explicit design requirement throughout this project

1.2. Problem Statement

The central problem this project addresses is that universities continue to identify struggling students too late and intervene too little, because their monitoring processes are reactive, manual, and unable to model the gradual way in which academic risk develops over time. Conventional end-of-semester reporting records failure after it has occurred, offering no meaningful window for corrective action, and even where early-warning tools exist, they typically produce isolated risk scores without the personalized guidance students need in order to act on them (Arnold & Pistilli, 2012).

In reality, the academic environment that institutions must manage has several characteristics that conventional approaches handle poorly:

- Student performance unfolds as a sequence over time, with engagement, submission behaviour, and assessment results evolving week by week rather than as a single fixed measurement.
- Students carry varying course loads, with one student enrolled in several subjects and another in only a few, so a monitoring system must be able to evaluate each active subject independently rather than assuming a uniform structure.
- The behavioural signals that precede academic failure, such as declining engagement combined with increasing submission delays, are non-linear and interact in ways that simple threshold rules cannot capture.
- Even when at-risk students are correctly identified, the personalized communication required to support them does not scale, because crafting individual messages for thousands of students manually is impractical.

Because of these mismatches, traditional monitoring approaches fall short in three critical ways. First, they detect risk only after it has materialized, eliminating the opportunity for timely intervention. Second, they treat performance as a static snapshot and therefore miss the temporal trajectory that reveals risk as it builds. Third, they stop at detection and place the entire burden of personalized response on already overstretched academic staff.

To address these shortcomings, this project proposes an intelligent, event-driven agent that models each student's per-subject academic trajectory using an LSTM network, converts the resulting risk scores into empathetic and actionable intervention messages through a language-model reasoning engine, and presents these to advising staff and students through an integrated dashboard. This approach offers a scalable path from passive observation toward proactive, personalized academic support.

1.3. Research Objectives

The primary objective of this research is to design, develop, and evaluate an AI-based agent that monitors student academic performance and enhances it through early, personalized intervention.

The specific objectives are:

1. To develop the predictive engine, by designing a workload-agnostic data preparation pipeline that transforms longitudinal student behavioural data into per-subject sequences and training an LSTM neural network that outputs a per-subject academic risk score between 0.0 and 1.0.
2. To implement the reasoning and delivery system, comprising a Large Language Model reasoning engine that converts the predicted risk scores into personalized, empathetic, subject-specific intervention messages, integrated through an event-driven Agent Controller, a MySQL data layer, and a web-based dashboard.
3. To evaluate and benchmark the system, by assessing its predictive performance using standard classification metrics such as AUC, accuracy, precision, and recall against baseline models such as Convolutional Neural Networks, and by assessing the quality of the generated intervention messages.

1.4. Scope of the study

This study focuses on the design and evaluation of an academic monitoring agent using the Open University Learning Analytics Dataset (OULAD) as its empirical foundation. The dataset provides longitudinal records of student demographics, registrations, virtual learning environment interactions, and assessment outcomes across multiple modules and presentations, making it well suited for developing and validating a sequential prediction model (Kuzilek et al., 2017).

Technically, the research is bounded by the use of an LSTM network as the predictive engine, a Large Language Model accessed through an external interface as the reasoning engine, and a system stack comprising Python, TensorFlow and Keras, MySQL, a Flask application programming interface, and a React user interface. The predictive task is framed as a binary classification of academic risk, where the model distinguishes enrollments likely to end in

withdrawal or failure from those likely to result in a pass or distinction, and produces a continuous risk score that is grouped into low, medium, and high tiers for intervention purposes.

The evaluation is centred on an early-warning checkpoint, using behavioural data available up to a configurable week of the presentation, in order to reflect the realistic conditions under which an institution would seek to act while the semester is still in progress. The study acknowledges that OULAD originates from a distance-learning institution, and it therefore treats the system as an architecture validated on online data, with engagement modelled as an abstracted input that can later be supplied by attendance or learning-management records in an in-person university. A full longitudinal study of intervention outcomes in a live institutional setting is beyond the scope of this project and is identified as future work.

1.5. Significance of the Study

This research is significant because it demonstrates a practical and scalable approach to a problem of genuine institutional importance: the timely identification and support of students at risk of academic failure. By modelling academic performance as a temporal sequence rather than a fixed snapshot, the study shows that deep learning can detect emerging risk earlier and more accurately than retrospective reporting allows, achieving strong predictive performance at an early-warning checkpoint well before a semester concludes (Santos & Henriques, 2024).

Beyond prediction, the study contributes a novel integration of sequential deep learning with language-model reasoning within a single autonomous agent. This pairing addresses a limitation that has constrained many earlier early-warning systems, namely that identifying at-risk students does little good if institutions lack the capacity to respond to each of them individually. By automating the generation of personalized, empathetic intervention messages, the system makes individualized academic support feasible at scale, building on evidence that large language models can produce feedback of a quality comparable to that of human instructors (Dai et al., 2023), and freeing advising staff to focus on the cases that most require human judgement.

The study is also significant in its attention to deployability and responsible design. By deliberately abstracting the engagement signal so that it can be supplied by either online interaction data or in-person attendance records, the architecture is positioned for adoption in universities that do not operate fully online, including institutions in the regional context. Furthermore, by explicitly excluding sensitive attributes from the predictive model and

documenting the associated fairness considerations, the study offers a template for building academic risk systems that are both effective and ethically grounded.

1.6. Definition of terms

To ensure clarity and consistent interpretation of the technical and economic concepts discussed in this study, the following terms are defined as follows:

Table 1.1 Operational Definition of Terms

Term	Definition
LSTM	Long Short-Term Memory is a recurrent neural network architecture capable of learning long-term dependencies in sequential data, making it suitable for modelling how student behaviour evolves over time.
Risk Score	A continuous value between 0.0 and 1.0 produced by the LSTM for each student-subject, where higher values indicate a greater likelihood of academic failure or withdrawal.
LLM	A Large Language Model is a deep learning model trained on extensive text data and capable of generating fluent, context-aware natural language, used here to produce personalized intervention messages.
OULAD	The Open University Learning Analytics Dataset, a publicly available dataset of student demographics, interactions, and assessment outcomes used to train and evaluate the model.
Agent Controller	The Python-based component that orchestrates the system by retrieving active subjects, running model inference, invoking the reasoning engine, and logging interventions.
Early Warning System	A monitoring system that identifies students at risk of academic failure early enough for corrective intervention to remain effective.
AUC	The Area Under the Receiver Operating Characteristic Curve, a metric that measures a classifier's ability to distinguish between at-risk and not-at-risk cases across all thresholds.
Enrollment	A single unit of analysis representing one student registered in one subject during one semester presentation.

CHAPTER TWO: LITERATURE REVIEW

2. Introduction to the Chapter

The paradigm of higher education infrastructure is undergoing a profound structural evolution, shifting away from retrospective administrative data logging toward predictive, real-time, human-centric computational ecosystems. Historically, institutional repositories served primarily as archival registries for academic transcript management, student enrollment records, and historical grade point averages. However, the contemporary expansion of distributed learning architectures, ubiquitous Virtual Learning Environments (VLEs), and complex institutional databases has created an unprecedented volume of granular, high-frequency student interaction data.

This chapter provides a rigorous, comprehensive, and multi-disciplinary critique of the existing body of empirical research underpinning these developments. It begins by mapping the conceptual trajectory of Educational Data Mining (EDM) and Learning Analytics (LA), tracking the transition from static, linear statistical heuristics to highly dynamic behavioral analysis models.

Following this, the chapter evaluates the mathematical and structural deployment of deep recurrent architectures—specifically Long Short-Term Memory (LSTM) networks—for multi-modal longitudinal student tracking, explicitly examining how these systems resolve the classic mathematical limitations of standard backpropagation through time.

The review then turns to the computational frontier of Generative AI, contextualizing Large Language Models (LLMs) not merely as static conversational interfaces, but as sophisticated, agentic reasoning layers capable of translating cold numerical risk metrics into highly contextualized, empathetic, and personalized pedagogical mentorship scripts.

Finally, this chapter establishes the ethical, legal, and sociotechnical frameworks necessary to safely govern these predictive architectures within modern university systems and concludes by delineating the precise research gaps this thesis aims to close.

2.1. Educational Data Mining (EDM) and Predictive Student Analytics

Educational Data Mining (EDM) focuses on developing methods for exploring unique types of data that come from educational settings to better understand students and the environments in which they learn (Romero and Ventura, 2010). The evolution of EDM is deeply intertwined with the digital transformation of higher education infrastructure. Historically, student performance prediction was an administrative exercise conducted post-facto, relying on static demographic indicators, socio-economic status, and cumulative historical Grade Point Averages (GPAs). Early frameworks primarily utilized traditional statistical methodologies, such as Logistic Regression, Support Vector Machines (SVM), Decision Trees, and Naive Bayes, to categorize students into binary "Pass" or "Fail" brackets at the end of an academic term (Asif et al., 2017).

However, traditional predictive models suffer from a fundamental structural limitation: they treat the academic semester as a static event rather than a dynamic, evolving process. Relying solely on static features introduces the "Cold Start Problem" in learning analytics, where the system can make generalizations based on past cohorts but fails to capture the immediate behavioral shifts of a student in a new course (Maldonado et al., 2021).

With the widespread adoption of Virtual Learning Environments (VLEs) like Moodle, Canvas, and Blackboard, universities now collect vast arrays of daily clickstream logs. These logs include quiz attempt frequencies, forum engagements, content access tracking, and resource download latencies. Research indicates that dynamic behavioral data provides a significantly more accurate reflection of student engagement, cognitive consistency, and academic self-regulation than demographic factors alone (Conijn et al., 2017).

A major challenge in early EDM research was the high reliance on manual feature engineering. Researchers had to aggregate daily clickstream data into semester-long averages or cumulative totals (e.g., "Total clicks per semester"). This static aggregation process effectively stripped away the temporal context, obscuring the behavioral variance of students. For instance, a model looking at cumulative clicks cannot differentiate between a student who consistently interacts with the VLE for two hours every week and a student who completely disengages for three months but crams 100 hours of activity into the final week before the exam (Gašević, Dawson and Rogers, 2016). Consequently, traditional classifiers struggled with early detection, often failing to yield actionable insights until the semester was too close to completion for effective institutional intervention.

2.2. Deep Learning and Longitudinal Time-Series Forecasting

To process the high-dimensional, sequential time-series data generated by modern VLEs without losing temporal context, learning analytics has shifted from traditional machine learning toward Deep Learning architectures. Standard feedforward neural networks fall short in this domain because they lack temporal memory; they treat a student's VLE interactions on Day 30 as completely independent of their behavior on Day 5.

Recurrent Neural Networks (RNNs) introduced recurrent connections to maintain an internal state across time-steps. However, standard RNNs suffer severely from the vanishing and exploding gradient problems when dealing with long academic semesters (e.g., sequences stretching over 200 days). This makes it mathematically impossible for the network to retain information from the early weeks of a semester by the time it processes the final weeks (Pascanu, Mikolov and Bengio, 2013).

Long Short-Term Memory (LSTM) networks resolve this structural limitation by introducing a specialized gating mechanism consisting of an input gate, a forget gate, and an output gate (Hochreiter and Schmidhuber, 1997). In the context of learning analytics, LSTMs can identify complex patterns of gradual academic disengagement, such as a subtle decline in weekly quiz scores coupled with a steady increase in assignment submission latency.

Table 2.1 Critical Comparison of Predictive Algorithms in Learning Analytics

Algorithm Class	Suitable Data Type	Temporal Awareness	Key Theoretical Limitation
Logistic Regression	Static (Demographics, GPA)	None	Cannot process time-varying weekly logs; assumes linear relationships.
Random Forest	Static / Tabular Aggregates	None	Desensitized to the chronological order of student actions.
Standard RNN	Short Sequential Data	Short-term memory	Suffers from vanishing gradients over long semantic intervals.
LSTM Networks	Multivariate Time-Series	Long-term memory	Requires large-scale, sequence-aligned datasets to avoid overfitting.

2.3. Generative AI and Automated Pedagogical Mentorship

While predictive modeling has achieved high statistical accuracy, a critical operational gap remains: **interpretability and communication**. Standard predictive engines produce quantitative risk outputs (e.g., a 0.72 probability of failure). For students, receiving a raw, unexplained risk percentage can induce unnecessary academic anxiety and learned helplessness. For university faculty, managing thousands of incoming numerical flags across a large student population creates an overwhelming cognitive and administrative load (Alam, 2021).

The emergence of Large Language Models (LLMs) and Generative AI presents a paradigm shift in student support systems. By deploying an agentic reasoning layer powered by Generative AI, a quantitative risk flag can be automatically translated into personalized, supportive, and context-aware feedback (Zawacki-Richter et al., 2019).

Modern agentic workflows utilize advanced Prompt Engineering techniques, such as Chain-of-Thought (CoT) prompting and role-specification. When an LLM is fed structured system context—combining a student's numerical risk score with specific behavioral deficiencies extracted from the VLE database (e.g., a 40% reduction in time spent on a specific module)—it can logically deduce the student's exact learning bottleneck. Rather than generating a generic alert, the agent drafts a highly targeted, empathetic intervention message that mimics a human mentor. This continuous automated feedback mechanism expands institutional capacity, ensuring that every student receives tailored guidance without increasing the administrative burden on human educators.

2.4. Ethical Considerations in AI-Driven Learning Analytics

The deployment of predictive engines and automated agents in higher education introduces significant ethical and legal challenges that must be addressed within any systemic design. The primary ethical concerns center around **data privacy, algorithmic bias, and student stigmatization** (Prinsloo and Slade, 2017).

1. **Data Privacy and Consent:** Learning analytics networks require access to sensitive student profiles, including demographic classifications and continuous behavioral tracking. System architectures must strictly comply with international data protection standards, such as the General Data Protection Regulation (GDPR), ensuring robust

data encryption at rest and in transit, anonymized data processing tokens, and transparent student opt-in protocols.

2. **Algorithmic Bias and Stigmatization:** If predictive models are heavily weighted on historical demographic data, they risk reinforcing institutional biases, unfairly flagging specific student groups based on background rather than current effort. Furthermore, labeling a student as "High Risk" early in a semester can lead to self-fulfilling prophecies or biased faculty perceptions. Therefore, contemporary research emphasizes that AI outputs should serve as a supportive "safety net" rather than an absolute academic judgment, and interventions should remain private and growth-oriented.

2.5. Identification of the Research Gap

Despite the independent advancements in predictive accuracy, sequence modeling, and generative communication, a structural disconnection persists within existing academic literature:

1. **Isolation of Educational AI Systems:** Current educational frameworks typically operate in silos. Predictive data engines generate risk metrics but lack automated, natural communication channels, while generative chatbots provide generic, conversational support without any grounding in real-time, empirical behavioral metrics.
2. **Fixed-Structure Workloads:** Existing predictive student monitoring systems are built around rigid, hard-coded course configurations. They assume every student maintains an identical course load (e.g., exactly 8 subjects per semester), making them unstable when deployed across diverse university ecosystems where individual credit hours and course loads fluctuate dynamically per student per semester.

This thesis addresses these gaps by engineering an integrated, **Closed-Loop Hybrid AI Framework**. By coupling a dual-input TensorFlow/Keras LSTM predictive engine with an asynchronous LLM reasoning controller over a dynamically scaling relational database, this system bridges the gap between early data detection and automated human-centric coaching, providing a scalable, workload-agnostic solution for modern higher education institutions.

CHAPTER 3: RESEARCH METHODOLOGY

3. Research Design

This study adopts an empirical, quantitative research design combined with an experimental system-engineering approach. The objective is to design, implement, and evaluate a predictive system capable of identifying academic risk while providing automated pedagogical feedback. The study leverages secondary data mapping to simulate real-time behavioral data streams.

By utilizing experimental evaluation, the system's predictive engine is tested under variable timelines to measure early-warning stability. The architectural performance is evaluated based on precision, recall, and receiver operating characteristic curves (ROC-AUC), ensuring that the quantitative findings are structurally validated.

3.1. System Development Methodology (CRISP-DM)

To engineer the system framework, this project adopts the **Cross-Industry Standard Process for Data Mining (CRISP-DM)** framework. CRISP-DM provides an iterative, structured lifecycle consisting of six discrete phases, ensuring that the machine learning engine aligns with real-world institutional objectives.

[Phase 1: Business Understanding] -> Defines institutional goals and intervention thresholds.

[Phase 2: Data Understanding] -> Explores and profiles the multivariate OULAD dataframes.

[Phase 3: Data Preparation] -> Handles padding, tensor shaping, and 3D matrix transformation.

[Phase 4: Modeling] -> Trains the multi-modal TensorFlow/Keras LSTM network.

[Phase 5: Evaluation] -> Audits predictive performance using early-warning milestones.

[Phase 6: Deployment] -> Integrates the Python engine with MySQL, Flask, and React.

- **Business/Institutional Understanding:** The primary objective is established: reducing student failure rates by identifying disengagement early. The operational constraint requires the system to maintain predictive accuracy within the first 30 days of a semester.
- **Data Understanding:** This phase involves profiling student behavioral footprints within virtual platforms, examining features such as login distributions, quiz submission frequencies, and interaction volumes.

- **Data Preparation:** Raw relational records are aggregated, cleaned, and transformed. This involves reshaping flat data into a **3D Tensor Matrix** suitable for deep sequential modeling.
- **Modeling:** The TensorFlow/Keras dual-input architecture is designed and trained. This phase utilizes sequence prefix cropping to expose the network to varying temporal windows.
- **Evaluation:** The predictive engine is audited using confusion matrices and recall metrics. This phase evaluates whether the model catches at least two-thirds of at-risk students.
- **Deployment:** The finalized parameters are serialized. The engine is then connected to a Flask API backend, a MySQL database, and a React user interface to close the loop.

3.2. Data Collection and Dataset Analysis

The framework is powered by the Open University Learning Analytics Dataset (OULAD), which contains information about 32,593 students, their demographic profiles, and their clickstream interactions within a VLE across 22 module presentations.

3.2.1 Relational Schema Analysis

The OULAD structure consists of seven interconnected dataframes that map the student learning lifecycle:

- **studentInfo:** Contains demographic profiles, including gender, region, education layer, age band, and disability markers.
- **studentRegistration:** Records the chronological date of student module registration and unregistration relative to the start of the term.
- **courses:** Defines the thematic modules (AAA to GGG) and their respective presentation timelines.
- **vle:** An inventory of all structural assets available within the virtual environment (e.g., quizzes, subpages, URLs, forums).
- **studentVle:** A massive longitudinal clickstream ledger containing daily interaction counts for every student across distinct assets.
- **assessments:** Defines the rules, dates, and relative weights of tutor-marked and computer-marked tests.

- **studentAssessment:** Records individual raw scores and submission dates for each assessment.

3.2.2 Applicability to Offline and Blended Universities

A common critique of learning analytics frameworks trained on datasets like OULAD is their applicability to traditional, offline brick-and-mortar universities. However, modern physical universities heavily utilize a blended learning model. Systems like Moodle, Canvas, and Google Classroom are standard extensions of physical lectures.

In a traditional university setting, this system can be deployed by routing daily activity logs from the campus LMS into the relational schema. Physical attendance registers can be encoded as a dynamic feature alongside quiz interactions, making the framework highly transferable to physical higher education environments.

3.3. System Requirements Analysis

To ensure structural scalability and runtime stability, the system requirements are defined across functional and non-functional specifications.

3.3.1 Functional Requirements

- **Automated Data Ingestion:** The system must periodically fetch new interaction logs from the transactional databases without manual intervention.
- **Continuous Risk Calculation:** The predictive engine must ingest the time-series vector and output a refined continuous failure probability score.
- **Threshold-Governed Policy Mapping:** The agent controller must categorize the score into distinct risk tiers (Low, Medium, High).
- **Generative Mentor Text Synthesis:** The system must pass high-risk context into an LLM to generate personalized, growth-oriented intervention text.
- **Asynchronous Dashboard Delivery:** The generated message must be rendered within a student-facing portal.

3.3.2 Non-Functional Requirements

- **Latency:** Inference pipelines and API communication must execute within 2 seconds.
- **Scalability:** The MySQL database and Flask routing layer must support concurrent queries for up to 30,000 active rows.

- **Security & Privacy:** Student identifiers must be cryptographically tokenized to ensure strict privacy compliance.

3.4. System Design and Architecture

The framework is built using a decoupled microservices architecture. This separates heavy tensor processing from user-facing frontend applications to maintain system performance.




Figure 3.1 system design and architecture

3.4.1 The Database Layer (MySQL)

The database stores student information across normalized tables (students, enrollments, daily_activity, agent_interventions), minimizing redundancy and ensuring efficient data retrieval.

3.4.2 The Inference Engine and Agent Controller

Written in Python and built on TensorFlow/Keras, this background service queries the daily activity logs, formats them into a 3D tensor, and loads lstm_risk_model.keras. If a student's risk profile crosses a set threshold, the controller formats a system prompt and triggers an API call to the LLM.

3.4.3 The Backend Application Interface (Flask API)

Flask serves as the RESTful coordination layer. It handles endpoint requests from the user interface, queries intervention states from MySQL, and passes data to the frontend in JSON format.

3.4.4 The Frontend Portal (React UI)

Built with React and styled with Tailwind CSS, this interface renders two distinct views:

- **The Student View:** Renders personalized, empathetic messages generated by the AI agent along with recommended study links.
- **The Advisor View:** Renders an analytics panel displaying risk trajectories and tracking intervention statuses.

3.5. Development Tools and Technologies

- **Programming Environment:** Python 3.11 for deep learning and data wrangling; JavaScript (ES6+) for user interface logic.
- **Deep Learning Framework:** TensorFlow 2.x with Keras for constructing the dual-input LSTM pathways.
- **Data Processing Libraries:** Pandas and NumPy for tensor transformations; Scikit-learn for stratified splitting.
- **Database Management System:** MySQL 8.0 for relational storage.
- **Backend Web Framework:** Flask 3.0 for API routing.
- **Frontend UI Library:** React 18 for rendering the responsive dashboard components.

3.6. Testing and Evaluation

The system validation pipeline splits the dataset into an **80/10/10 Stratified Split**, ensuring equal class distributions across the train, validation, and test sets.

- **Predictive Metrics:** Evaluated using Macro F1-Score, ROC-AUC, and Class-Specific Recall to maximize the detection of at-risk students.
- **System Integration Testing:** Validated by executing end-to-end dry runs: injecting raw values into MySQL, running inference.py, generating the LLM text, and verifying the interface display.

3.7. Ethical Considerations

To safeguard student interests, the framework incorporates two primary ethical guardrails:

1. **Data Anonymization:** Student institutional IDs are tokenized before model ingestion, ensuring that personal identifying information is never exposed to external APIs.

2. **Growth-Oriented Interventions:** The LLM prompt rules strictly prohibit alarmist or punitive language. The system functions exclusively as a supportive "safety net" to encourage re-engagement.

3.8. Schedule Feasibility

The development timeline follows a structured sequence over a six-month window, organized to ensure systematic validation:

Table 3.1 Project Implementation Timeline

Phase	Focus Area	Key Deliverable	Duration
Phase 1	Preprocessing & Wrangling	3D Tensor generation and .npz serialization.	Weeks 1–4
Phase 2	Model Development	Training TensorFlow/Keras dual-input LSTM with weighted loss.	Weeks 5–10
Phase 3	Integration Layer	Engineering AgentController.py and Flask API.	Weeks 11–14
Phase 4	Frontend Engineering	Building React interface components and dashboards.	Weeks 15–18
Phase 5	Evaluation & Debugging	Running integration testing and generating ROC-AUC metrics.	Weeks 19–22
Phase 6	Documentation	Finalizing the thesis manuscript and system analysis.	Weeks 23–24

CHAPTER FOUR: RESULTS AND DISCUSSIONS

4. Introduction

This chapter presents the empirical findings, quantitative performance metrics, and analytical evaluations of the trained dual-input Long Short-Term Memory (LSTM) student academic performance predictive network. The primary focus is to analyze the model's performance on a held-out test partition, examine the implications of the intentional precision-recall trade-off, and critique the longitudinal stability of the system using an early-warning milestone analysis. Finally, the chapter details the live runtime software deployment verification, evaluating the operational user interfaces engineered for both institutional academic advisors and individual end-user students.

4.1. Architectural Convergence Analysis

Before evaluating the test metrics, it is essential to analyze the structural convergence of the optimization pipeline to ensure the model generalized effectively without overfitting. The network was trained using the Adam optimizer with binary cross-entropy loss adjusted by balanced class weights, with a minority-class (Fail/Withdrawn) weight of \$1.09\$ to handle the target imbalance.

The optimization process monitored validation loss over a maximum of 50 epochs, protected by an early stopping callback (monitoring validation AUC) with a patience of 5 epochs. The system achieved optimal generalization and structural stability at **Epoch 31**, with the training loop programmatically terminating at **Epoch 39**.

The loss trajectory demonstrates an ideal learning curve, where both curves decline sharply within the first 25 epochs and stabilize asymptotically. The minor, consistent gap between the final training loss (\$0.89386\$) and validation loss (\$0.88868\$) confirms that the integration of dropout layers ($p = 0.3$) successfully prevented overfitting, establishing a robust foundation for test-set evaluation.

4.2. Model Performance and Validation on the Test Set

The finalized model parameters, serialized within `lstm_risk_model.keras`, were evaluated using a completely held-out test partition consisting of **4,238 active student profiles** monitored over their full sequence lengths.

4.2.1 Area Under the Curve (ROC-AUC) and General Accuracy

The predictive engine achieved a **Receiver Operating Characteristic Area Under the Curve (ROC-AUC) score of 0.9839**. This score indicates a statistically stable capability to differentiate between passing students and those displaying failure-prone behavioral trajectories. The overall raw classification accuracy settled at **94%**. This high score indicates a strong capability to differentiate between passing students and those displaying failure-prone behavioral trajectories. The overall classification accuracy of 94% confirms that, given a student's complete behavioural sequence, the model separates the two outcome classes with high reliability. As discussed in the early-warning milestone analysis below, this discriminative power is strongest once a substantial portion of the semester has elapsed and is more modest at earlier checkpoints.

4.2.2 Confusion Matrix Evaluation

To understand the exact distribution of classification boundaries, the model's test predictions were analyzed using a standard classification threshold ($\theta = 0.5$). The resulting confusion matrix maps the alignment between actual outcomes and model predictions:

Table 4.1 Test Set Confusion Matrix Boundaries ($N = 4,238$)

	Predicted Low Risk (Pass/Distinction)	Predicted High Risk (Academic Failure)
Actual Low Risk (Pass/Distinction)	2,222 (True Negatives)	91 (False Positives)
Actual High Risk (Academic Failure)	167 (False Negatives)	1,758 (True Positives)

4.2.3 Detailed Classification Metrics

To further evaluate the framework's effectiveness as an institutional intervention tool, class-specific precision, recall, and F1-scores were computed across the test partition:

Table 4.2 Comprehensive Test Set Classification Report

Target Class Label	Precision	Recall (Sensitivity)	F1-Score	Support Base
Class 0.0 (Pass / Distinction)	0.93	0.96	0.95	2,313

Class	1.0	(Academic	0.95	0.91	0.93	1,925
Failure)						
Macro Average Summary			0.94	0.94	0.94	4,238

4.3. Discussion of the Precision-Recall Asymmetry

A critical analysis of the classification report highlights a strong and well-balanced detection profile: **the Fail/Withdrawn class achieves a high Recall of 0.91 together with a high Precision of 0.95.** In machine learning pipelines engineered for educational institutions, this specific balance is highly desirable.

The high failure recall proves that **the system successfully catches 91% of all students who ultimately fail or withdraw from the course**, mapping them into active intervention pipelines. At the same time precision remains high, with only 91 False Positives (students flagged as at-risk who in fact passed) across the 4,238 test enrollments.

From a pedagogical perspective, the cost of a False Positive is remarkably low—it simply triggers an automated, supportive reach-out message or an academic tutoring nudge through the student dashboard. Conversely, the cost of a False Negative (missing a student genuinely heading toward failure) is catastrophic, leading to student dropout or course failure. Therefore, the empirical results validate that the use of `balanced class weights` successfully constructed a model that is both highly sensitive and precise on the full-sequence task.

4.4. Early-Warning Milestone Analysis

To demonstrate that the model operates as a proactive, real-time tracking engine rather than a post-facto classifier, sequence prefix cropping was applied to truncate the test sequences at distinct chronological milestones throughout the semester.

Table 4.3 Longitudinal Milestone Analysis Performance Over Time

Milestone	Academic	Model	Failure	Catch	F1-Score
Horizon	Semester Phase	ROC-AUC	Rate (Recall)		Summary
Day 30	Early-Term	0.7030	35.58%		0.4663
	Boundary				

Day 60	Mid-Term Checkpoint 1	0.8111	36.16%	0.5073
Day 90	Mid-Term Checkpoint 2	0.8368	44.52%	0.5866
Day 120	Late-Term Transition	0.9115	58.34%	0.7236
Day 180	Pre-Examination Anchor	0.9575	78.29%	0.8644
Day 240	Final Semester Phase	0.9808	90.23%	0.9254

The milestone analysis provides the clearest evidence of how early the system becomes useful. Rather than being constant, the model’s predictive performance improves steadily as more of the semester is observed, rising from an AUC of 0.70 at Day 30 to 0.98 once the full sequence is available.

One month into the academic term (Day 30) the model already captures **35.6%** of the eventual failure profiles with an ROC-AUC of **0.70**. By the week-13 early-warning checkpoint this rises to an AUC of 0.84 (44.5% recall), and it continues to climb to 0.91 by Day 120 and 0.98 once the full sequence is available. This trajectory shows that meaningful early detection is possible within the first weeks, while the strongest discrimination is naturally achieved later in the term as more behavioural evidence accumulates. The sequence prefix cropping strategy confirms that the LSTM extracts useful disengagement signals from incomplete interaction windows, with recall improving as the observation window lengthens.

4.5. System Verification and Dashboard Interface Analysis

To validate the structural execution of the predictive model when integrated into the full web application architecture, the system runtime performance was inspected via user interface outputs mapped back to the relational MySQL database.

4.5.1 Academic Advisor Analytics Dashboard View

The administrative portal serves as the primary oversight layer for institutional educators to identify at-risk trends across student cohorts.




Figure 4.1 Institutional Advisor Analytics Console and Cohort Risk Manifest

As illustrated in **Figure 4.1**, the advisor dashboard translates background student attributes and current semester activity records directly into dynamic status tags. Based on institutional threshold values, students crossing specific probability bounds are automatically assigned risk tags, allowing tutors to segment large-scale classrooms into priority intervention clusters instantly.

4.5.2 Advisor In-Depth Profile Review

When an advisor selects a flagged profile, the interface reads directly from the Flask API endpoint to generate an in-depth demographic and interaction assessment overview.




Figure 4.2 Student #6516 Single Profile Advisor Analysis Screen

The performance tracking for **Student #6516** (Course Module: AAA, Presentation: 2014J) demonstrates the core utility of the system under live verification (**Figure 4.3**). The execution engine automatically assesses student history up to the latest operational activity timestamp (auto-detected at Day 4).

The underlying TensorFlow/Keras model processes the multi-modal input vectors to register a failure probability of **22.9822%**. Although a 22.98% score represents stable baseline progress over full-semester projections, the early-term application triggers an alert policy due to short sequence click drift, recommending early follow-up.

4.5.3 End-User Student Portal Messaging Layer

The downstream reasoning layer closes the loop by outputting conversational feedback directly to the student portal interface, driven asynchronously by the Large Language Model reasoning class.




Figure 4.3 Student Dashboard

As shown in **Figure 4.3**, the student-facing React portal displays the output generated by the AI Mentor Agent based on database interaction logs. Rather than exposing the student to clinical risk scores or alarmist metrics, the dashboard renders a highly tailored message.

The agent highlights specific learning tasks, such as re-engaging with the digital quiz modules, while embedding actionable links to library and university support ecosystems. This layout demonstrates that the closed-loop system successfully bridges the gap between deep sequential failure prediction and empathetic academic support.

4.6. Comparison with Existing Systems

To situate the proposed system within the wider learning-analytics literature, Table 4.4 compares its predictive performance and functional scope against representative early-warning systems trained on the OULAD or comparable learning-management data. The comparison

considers not only raw predictive metrics but also the checkpoint at which prediction is made and, critically, whether each system closes the loop by generating automated interventions.

Table 4.4 Comparison of the Proposed System with Existing Early-Warning Systems

System / Study	Model	Dataset	Reported performance	Automated intervention?
Aljohani et al. (2019)	LSTM	OULAD (32,593 students)	Acc 95.2%, Prec 93.5%, Recall 75.8% (full-course data)	No — risk flag only
Liu et al. (2023)	LSTM (weekly clicks)	OULAD (5,341 students)	Acc 89.3%, F1 92.7%, AUC 0.913	No — risk flag only
Santos & Henriques (2024)	LSTM / classical ML	LMS clickstream logs	AUC ≈ 0.76 (best ML); LSTM lower but viable earlier	No — risk flag only
Course Signals – Arnold & Pistilli (2012)	Logistic regression	Purdue SIS / LMS	Improved retention; traffic-light flag	No — flag, manual outreach
Proposed system (this study)	Dual-input LSTM + LLM	OULAD (28,297 enrollments)	AUC 0.984, Acc 94%, Recall 91% (full seq); AUC 0.84 at week 13	Yes — LLM-generated message

As Table 4.4 shows, on the full behavioural sequence the proposed LSTM (AUC 0.984) is competitive with, and in several cases exceeds, the strongest OULAD-based systems reported in the literature, such as Liu et al. (2023) at AUC 0.913 and Aljohani et al. (2019) at 95% accuracy. A fair reading must note that these studies differ in their exact task definitions, cohorts, and evaluation points, so the comparison is indicative rather than a controlled head-to-head. At the earlier week-13 checkpoint the proposed model is necessarily more modest (AUC 0.84), which is the realistic operating point for genuinely early intervention. The more important distinction, however, is functional: every comparison system is prediction-only, terminating at a risk flag and delegating the design of the intervention to human staff.

The defining contribution of this study is therefore not a marginal gain in discrimination but the integration of a Large Language Model reasoning layer that converts each risk score into a personalised, empathetic intervention message. This closed-loop capability, combined with a workload-agnostic data pipeline that adapts to varying course loads, is absent from all of the compared systems and represents the principal advance offered by the proposed architecture.

4.7. Comparative Model Evaluation

To establish that the chosen sequential architecture is justified, the proposed LSTM was benchmarked against three alternative models trained and evaluated on the identical student-level split: a one-dimensional Convolutional Neural Network (CNN) with the same dual-input structure, a Random Forest, and a Logistic Regression classifier. For the two classical models

the weekly sequence was flattened into a single feature vector and concatenated with the static features. Table 4.5 reports the results on the held-out test set, and Figure 4.4 visualises them.

Table 4.5 Comparative Performance of the LSTM, CNN, and Classical Baselines

Model	AUC	Accuracy	Precision	Recall	F1
LSTM (proposed)	0.984	0.939	0.951	0.913	0.932
CNN (1D conv)	0.977	0.927	0.956	0.878	0.916
Random Forest	0.980	0.937	0.968	0.890	0.927
Logistic Regression	0.980	0.933	0.946	0.905	0.925




Figure 4.4 Model comparison across evaluation metrics on the held-out test set

All four models perform strongly on the full-sequence task, which reflects the high separability of pass and fail outcomes once a complete behavioural record is available. The proposed LSTM nonetheless attains the highest AUC (0.984), narrowly ahead of the Random Forest and Logistic Regression (both 0.980) and the CNN (0.977), and it achieves the best recall among the deep models (0.913). The narrowness of these margins is itself informative: it indicates that most of the predictive signal is captured by the engineered weekly features, while the recurrent architecture provides a small but consistent advantage in modelling their temporal ordering. Crucially, the LSTM also produces a smoothly increasing risk trajectory over the semester (Section 4.4), a property the flattened classical models cannot offer and which is essential for the early-warning function of the system.

CHAPTER FIVE: CONCLUSIONS AND RECOMMENDATIONS

5. Introduction

This final chapter synthesizes the empirical outcomes, architectural milestones, and pedagogical insights generated throughout the development and evaluation of the Closed-Loop Hybrid AI Framework. Having successfully engineered a multi-modal predictive pipeline and integrated it with an agentic large language model reasoning layer, this section translates those computational results into definitive academic conclusions. Furthermore, it outlines actionable, growth-oriented recommendations for higher education administrative bodies and charts potential trajectories for future advancements in adaptive learning analytics frameworks.

5.1. Summary of Empirical Findings

The quantitative and software engineering investigations executed across this study yielded critical insights into the capability of advanced deep learning pipelines to mitigate student retention risks:

- **Predictive Performance Foundations:** The multi-modal TensorFlow/Keras model achieved a stable Receiver Operating Characteristic Area Under the Curve (ROC-AUC) score of 0.9839 and an overall test accuracy of 94% across a held-out test cohort of 4,238 student profiles, evaluated on the full behavioural sequence.
- **Sensitivity Optimization Success:** By training with balanced class weights computed from the data (the cohort is approximately 46% at-risk, a near-balanced distribution rather than a severe imbalance), the model learns the Fail/Withdrawn class effectively, achieving a failure-class recall of 91% together with a precision of 95% on the full-sequence test set.
- **Early-Warning Longitudinal Stability:** Sequence prefix cropping showed that predictive performance improves steadily as the semester progresses, rising from an AUC of 0.70 at Day 30 to 0.84 at the week-13 early-warning checkpoint and 0.98 on the full sequence. Early detection is therefore feasible within the first weeks, while the strongest discrimination is achieved later in the term, confirming that institutions can begin risk monitoring well before full-semester aggregates are available.
- **Closed-Loop System Integration:** Dynamic runtime validation verified that the combination of a React frontend, Flask backend API, and a normalized MySQL

repository successfully executes automated, continuous end-to-end processing. The architecture accurately evaluates student trajectories (such as Student #6516 reflecting a 22.98% risk on Day 4) and synthesizes empathetic, context-aware pedagogical guidance using Large Language Models without increasing manual faculty monitoring workloads.

5.2. Conclusions

The technical execution of this research provides a definitive solution to the structural limitations of traditional, reactive student retention monitoring. Based on the comparative performance and milestone evaluations, this study concludes that fusing multivariate sequential deep learning (LSTM) with demographic contextual embeddings (MLP) establishes a robust framework for continuous academic forecasting. Traditional linear approaches and tabular machine learning tools require macro-level aggregate feature totals, meaning they fail to capture intra-semester behavioral volatility and patterns of gradual disengagement until the term is too close to completion for effective outreach.

Furthermore, this study demonstrates that bridging the gap between raw statistical numbers and compassionate pedagogical action requires an automated reasoning layer. Leaving raw risk probabilities as numerical metrics leads to faculty alert fatigue and induces severe academic anxiety within vulnerable students. By implementing an agentic large language model architecture governed by conditional threshold matrices, institutions can convert complex high-dimensional feature anomalies into natural, supportive, and growth-oriented feedback loop notifications.

Crucially, because the underlying data extraction layers were engineered to process daily individual transactional interaction tokens dynamically, the system stands as a reliable, workload-agnostic monitoring framework capable of adapting seamlessly to the flexible academic course loads of modern higher education.

5.3. Recommendations

Based on the empirical findings, quantitative milestone performance indicators, and closed-loop system engineering frameworks validated throughout this study, the following comprehensive recommendations are proposed for academic, administrative, and technological stakeholders in higher education institutions:

Higher education administrative bodies should formally abandon late-stage, retrospective end-of-semester grade reporting as their primary monitoring mechanism. Instead, institutions must mandate the integration of dynamic, continuous early-warning pipelines to detect student disengagement features within the first 30 days of an academic term.

To support this shift, academic policies must ensure that automated risk scores are never utilized as high-stakes punitive metrics, such as for automated course disqualification or scholarship cancellations. Tutors and advisors must handle risk metrics with absolute confidentiality, treating the system outputs strictly as diagnostic triggers for growth-oriented outreach to protect vulnerable student demographics from algorithmic labeling and learned helplessness.

Universities should reorganize their tutoring models to align with the risk threshold matrices defined by automated systems. This enables advisors to save time on low-risk students and focus on scheduling personalized meetings for high-risk students, ensuring human intervention is directed exactly where it is needed most.

From a software engineering perspective, university IT departments and system architects should avoid monolithic designs and instead deploy decoupled, microservice-based multi-tier systems. Utilizing an architecture similar to the React frontend, Flask backend API, and normalized MySQL pipeline compiled in this thesis significantly minimizes runtime latency and ensures optimal real-time performance.

Furthermore, institutional relational databases must be structurally optimized to support continuous, event-driven data streaming from virtual learning environments. Clickstream logs—including forum activity, quiz access pacing, and document downloading latency indicators—should automatically pipe into raw transactional tables on a daily schedule to maintain prediction accuracy.

Finally, to expand the framework's versatility across traditional brick-and-mortar and blended campuses, the database schema should expand to ingest physical classroom metadata. Fusing features such as campus turnstile access logs, physical laboratory attendance sheets, smart classroom interactions, and digital library asset checkouts into the predictive pipeline will optimize accuracy across hybrid educational environments.

5.4. Limitations and Future Work

While the framework successfully demonstrated early risk identification and closed-loop messaging capabilities, several constraints outline boundaries for future machine learning research:

1. **Context-Insensitive Baseline Thresholds:** The risk categories are currently regulated by fixed static policy thresholds (e.g., 30% and 60%). Future iterations should integrate adaptive thresholding, where the system action boundaries are dynamically calibrated based on specific course difficulty metrics and structural module drop rates.
2. **Expansion of Sequential Architectures:** While the dual-input LSTM network effectively mitigated the vanishing gradient crisis across longitudinal student histories, future empirical setups should evaluate the comparative performance of **Transformer-Based Attention Architectures** (such as TabTransformer or behavioral BERT encodings) to analyze if self-attention mechanisms compress high-dimensional clickstream distributions with lower training overhead.
3. **Multi-Agent Evaluation Loops:** The system currently relies on a single LLM advisory node to draft intervention messages. Future software development should explore multi-agent architectures, where an auxiliary "Critic Agent" audits the synthesized mentor text against historical student interaction profiles to optimize the emotional resonance and readability metrics of every notification before final dispatch.

REFERENCE

- Arnold, K. E., & Pistilli, M. D. (2012). Course Signals at Purdue: Using learning analytics to increase student success. In *Proceedings of the 2nd International Conference on Learning Analytics and Knowledge* (pp. 267–270). ACM. <https://doi.org/10.1145/2330601.2330666>
- Dai, W., Lin, J., Jin, H., Li, T., Tsai, Y.-S., Gašević, D., & Chen, G. (2023). Can large language models provide feedback to students? A case study on ChatGPT. In *2023 IEEE International Conference on Advanced Learning Technologies (ICALT)* (pp. 323–325). IEEE. <https://doi.org/10.1109/ICALT58122.2023.00100>
- Ghazvini, A., Sharef, N. M., & Sidi, F. B. (2024). Prediction of course grades in computer science higher education program via a combination of loss functions in LSTM model. *IEEE Access*, 12, 30220–30241. <https://doi.org/10.1109/ACCESS.2024.3351186>
- Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural Computation*, 9(8), 1735–1780. <https://doi.org/10.1162/neco.1997.9.8.1735>
- Kuzilek, J., Hlosta, M., & Zdrahal, Z. (2017). Open University Learning Analytics dataset. *Scientific Data*, 4, Article 170171. <https://doi.org/10.1038/sdata.2017.171>
- Niang, M., et al. (2026). Machine learning and deep learning for dropout prediction in higher education: A review. *Computers*, 15(3), Article 164. <https://doi.org/10.3390/computers15030164>
- Romero, C., & Ventura, S. (2020). Educational data mining and learning analytics: An updated survey. *WIREs Data Mining and Knowledge Discovery*, 10(3), Article e1355. <https://doi.org/10.1002/widm.1355>
- Santos, R., & Henriques, R. (2024). Using early clickstream data to identify at-risk students in higher education: An LSTM-based approach. In *EDULEARN24 Proceedings* (pp. 6216–6225). IATED. <https://doi.org/10.21125/edulearn.2024.1471>

- Asif, R., Merceron, A., Ali, S.A. and Haider, N.G., 2017. Analyzing undergraduate students' performance using educational data mining. *Computers & Education*, 113, pp.177-194.
- Conijn, R., Snijders, C., Kleingeld, A. and Matzat, U., 2017. Predicting student performance from LMS data: A comparison of 17 blended courses. *Journal of Educational Technology & Society*, 20(3), pp.249-260.
- Hussain, M., Zhu, W., Zhang, W. and Abidi, S.M.R., 2018. Student engagement predictions in open university learning analytics dataset using deep learning. *IEEE Access*, 6, pp.20780-20792.
- Maldonado, J.J., Han, J.A., Khan, I.O. and Boroowa, A., 2021. Mining early indicators of student success from clickstream data in virtual learning environments. *Computers in Human Behavior*, 122, p.106843.
- Romero, C. and Ventura, S., 2010. Educational data mining: A review of the state of the art. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 40(6), pp.601-618.